

---

# REAL ESTATE

---

Project1




JULY 6, 2021  
rayasalehalshammri@gmail.com

## Problem Statement:

- A banking institution requires actionable insights into mortgage-backed securities, geographic business investment, and real estate analysis.
- The mortgage bank would like to identify potential monthly mortgage expenses for each region based on monthly family income and rental of the real estate.
- A statistical model needs to be created to predict the potential demand in dollars' amount of loan for each of the region in the USA. Also, there is a need to create a dashboard which would refresh periodically post data retrieval from the agencies.
- The dashboard must demonstrate relationships and trends for the key metrics as follows: number of loans, average rental income, monthly mortgage and owner's cost, family income vs mortgage cost comparison across different regions. The metrics described here do not limit the dashboard to these few.

## Dashboard image:



## Dashboard link:

[https://public.tableau.com/views/RealEstateProject-1\\_16255218497970/RealEstateDashboard?:language=en-US&:display\\_count=n&:origin=viz\\_share\\_link](https://public.tableau.com/views/RealEstateProject-1_16255218497970/RealEstateDashboard?:language=en-US&:display_count=n&:origin=viz_share_link)

## project in python link:

<https://github.com/rayas711/Real-Estate-Project1/blob/master/real-estate-project1-rayा.ipynb>

## Output image:

```
In [1]: import time
import random
from math import *
import operator
import pandas as pd
import numpy as np

In [2]: #Plotting Libraries
import matplotlib
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
%matplotlib inline

import seaborn as sns
sns.set(style="white", color_codes=True)
sns.set(font_scale=1.5)

Week1, 1- Import data:

In [3]: df_train=pd.read_csv("../input/realestateproject1/train.csv")
df_test=pd.read_csv("../input/realestateproject1/test.csv")

In [4]: df_train.columns
Out[4]: Index(['UID', 'BLOCKID', 'SUMLEVEL', 'COUNTYID', 'STATEID', 'state',
       'state_ab', 'city', 'place', 'type', 'primary', 'zip_code', 'area_code',
       'lat', 'lng', 'Aland', 'Awater', 'pop', 'male_pop', 'female_pop',
       'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight',
       'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25',
       'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50',
       'univise_samples', 'used_samples', 'hi_mean', 'hi_median', 'hi_stdev',
       'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median',
       'family_stdev', 'family_sample_weight', 'family_samples',
       'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',
       'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean',
       'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
       'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
       'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
       'hs_degree_male', 'hs_degree_female', 'male_age_mean',
       'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
       'male_age_samples', 'female_age_mean', 'female_age_median',
       'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
       'pct_own', 'married', 'married_snp', 'separated', 'divorced'],
      dtype='object')

In [5]: df_test.columns
Out[5]: Index(['UID', 'BLOCKID', 'SUMLEVEL', 'COUNTYID', 'STATEID', 'state',
       'state_ab', 'city', 'place', 'type', 'primary', 'zip_code', 'area_code',
       'lat', 'lng', 'Aland', 'Awater', 'pop', 'male_pop', 'female_pop',
       'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight',
       'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25',
       'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50',
       'univise_samples', 'used_samples', 'hi_mean', 'hi_median', 'hi_stdev',
       'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median',
       'family_stdev', 'family_sample_weight', 'family_samples',
       'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',
       'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean',
       'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
       'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
       'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
       'hs_degree_male', 'hs_degree_female', 'male_age_mean',
       'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
       'male_age_samples', 'female_age_mean', 'female_age_median',
       'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
       'pct_own', 'married', 'married_snp', 'separated', 'divorced'],
      dtype='object')

In [6]: len(df_train)
Out[6]: 27921

In [7]: len(df_test)
Out[7]: 11709

In [8]: df_train.head()
Out[8]:
   UID BLOCKID SUMLEVEL COUNTYID STATEID state state_ab city place type ... female_age_mean female_age_median female_z
0 267822 NaN 140 53 36 New York NY Hamilton Hamilton City ... 44.48629 45.33333
1 246444 NaN 140 141 18 Indiana IN South Bend Roseland City ... 36.48391 37.58333
2 245683 NaN 140 63 18 Indiana IN Danville Danville City ... 42.15810 42.83333
3 279653 NaN 140 127 72 Puerto Rico PR San Juan Guaynabo Urban ... 47.77526 50.58333
4 247218 NaN 140 161 20 Kansas KS Manhattan Manhattan City City ... 24.17893 21.58333
5 rows × 80 columns
```

```
In [9]: df_test.head()
Out[9]:
   UID BLOCKID SUMLEVEL COUNTYID STATEID state state_ab city place type ... female_age_mean female_age_median fe
0 255504 NaN 140 163 26 Michigan MI Detroit Dearborn Heights City CDP ... 34.78682 33.75000
1 252076 NaN 140 1 23 Maine ME Auburn Auburn City City ... 44.23451 46.66667
2 276314 NaN 140 15 42 Pennsylvania PA Pine City Monticello Monticello Borough ... 41.62426 44.50000
3 248814 NaN 140 231 21 Kentucky KY Monticello Monticello City City ... 44.81200 48.00000
4 286865 NaN 140 355 48 Texas TX Corpus Christi Edroy Town ... 40.66618 42.66667
```

```
In [10]: df_train.describe()
```

```
Out[10]:
```

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	zip_code	area_code	lat	lng	ALand	... femal
count	27321.000000	0.0	27321.0	27321.000000	27321.000000	27321.000000	27321.000000	27321.000000	27321.000000	27321.000000	2.732100e+04 ... ;
mean	257331.996303	NaN	140.0	85.646426	28.271805	50081.999524	596.507668	37.508813	-91.283394	1.295106e+08	...
std	21343.859725	NaN	0.0	98.333097	16.392846	29558.115660	232.497482	5.588268	16.343816	1.275531e+09	...
min	220342.000000	NaN	140.0	1.000000	1.000000	602.000000	201.000000	17.929085	-165.453872	4.113400e+04	...
25%	238816.000000	NaN	140.0	29.000000	13.000000	26554.000000	405.000000	33.889064	-97.816067	1.799408e+06	...
50%	257220.000000	NaN	140.0	63.000000	28.000000	47715.000000	614.000000	38.755183	-86.554374	4.8666940e+06	...
75%	275818.000000	NaN	140.0	109.000000	42.000000	77093.000000	801.000000	41.380606	-79.782503	3.359820e+07	...
max	294334.000000	NaN	140.0	840.000000	72.000000	99925.000000	989.000000	67.074018	-65.379332	1.039510e+11	...

8 rows × 74 columns

```
In [11]: df_test.describe()
```

```
Out[11]:
```

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	zip_code	area_code	lat	lng	ALand	... female
count	11709.000000	0.0	11709.0	11709.000000	11709.000000	11709.000000	11709.000000	11709.000000	11709.000000	1.170900e+04	1 ... ;
mean	257525.004783	NaN	140.0	85.710650	28.489198	50123.418398	593.598514	37.405491	-91.340228	1.095500e+08	...
std	21466.372658	NaN	0.0	99.304334	16.607262	29775.134038	232.074263	5.625904	16.407818	7.624940e+08	...
min	220336.000000	NaN	140.0	1.000000	1.000000	601.000000	201.000000	17.965835	-166.770797	8.299000e+03	...
25%	238819.000000	NaN	140.0	29.000000	13.000000	25570.000000	404.000000	33.919813	-97.816561	1.718680e+06	...
50%	257651.000000	NaN	140.0	61.000000	28.000000	47362.000000	612.000000	38.618092	-86.643344	4.835000e+06	...
75%	278300.000000	NaN	140.0	109.000000	42.000000	77406.000000	787.000000	41.232973	-79.697311	3.204540e+07	...
max	294333.000000	NaN	140.0	810.000000	72.000000	99929.000000	989.000000	64.804269	-65.695344	5.520166e+10	...

8 rows × 74 columns

```
In [12]: df_train.info()
```

```
In [12]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27321 entries, 0 to 27320
Data columns (total 80 columns):
 #   Column          Non-Null Count  Dtype  
--- 
 0   UID             27321 non-null   int64  
 1   BLOCKID         0 non-null     float64 
 2   SUMLEVEL        27321 non-null   int64  
 3   COUNTYID        27321 non-null   int64  
 4   STATEID         27321 non-null   int64  
 5   state            27321 non-null   object  
 6   state_ab        27321 non-null   object  
 7   city             27321 non-null   object  
 8   place            27321 non-null   object  
 9   type             27321 non-null   object  
 10  primary          27321 non-null   object  
 11  zip_code         27321 non-null   int64  
 12  area_code        27321 non-null   int64  
 13  lat              27321 non-null   float64 
 14  lng              27321 non-null   float64 
 15  ALand            27321 non-null   int64  
 16  Awater           27321 non-null   int64  
 17  pop              27321 non-null   int64  
 18  male_pop         27321 non-null   int64  
 19  female_pop       27007 non-null   float64 
 20  rent_mean        27007 non-null   float64 
 21  rent_median      27007 non-null   float64 
 22  rent_stdev       27007 non-null   float64 
 23  rent_sample_weight 27007 non-null   float64 
 24  rent_samples      27007 non-null   float64 
 25  rent_gt_10        27007 non-null   float64 
 26  rent_gt_15        27007 non-null   float64 
 27  rent_gt_20        27007 non-null   float64 
 28  rent_gt_25        27007 non-null   float64 
 29  rent_gt_30        27007 non-null   float64 
 30  rent_gt_35        27007 non-null   float64 
 31  rent_gt_40        27007 non-null   float64 
 32  rent_gt_50        27007 non-null   float64 
 33  universe_samples  27321 non-null   int64  
 34  used_samples      27321 non-null   int64  
 35  hi_mean           27053 non-null   float64 
 36  hi_median          27053 non-null   float64 
 37  hi_stdev           27053 non-null   float64
```

```
In [13]: df_test.info()
```

```
In [13]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11709 entries, 0 to 11708
Data columns (total 80 columns):
 #   Column          Non-Null Count  Dtype  
--- 
 0   UID             11709 non-null   int64  
 1   BLOCKID         0 non-null     float64 
 2   SUMLEVEL        11709 non-null   int64  
 3   COUNTYID        11709 non-null   int64  
 4   STATEID         11709 non-null   int64  
 5   state            11709 non-null   object  
 6   state_ab        11709 non-null   object  
 7   city             11709 non-null   object  
 8   place            11709 non-null   object  
 9   type             11709 non-null   object  
 10  primary          11709 non-null   object  
 11  zip_code         11709 non-null   int64  
 12  area_code        11709 non-null   int64  
 13  lat              11709 non-null   float64 
 14  lng              11709 non-null   float64 
 15  ALand            11709 non-null   int64  
 16  Awater           11709 non-null   int64  
 17  pop              11709 non-null   int64  
 18  male_pop         11709 non-null   int64  
 19  female_pop       11709 non-null   int64  
 20  rent_mean        11561 non-null   float64 
 21  rent_median      11561 non-null   float64 
 22  rent_stdev       11561 non-null   float64 
 23  rent_sample_weight 11561 non-null   float64 
 24  rent_samples      11561 non-null   float64 
 25  rent_gt_10        11560 non-null   float64 
 26  rent_gt_15        11560 non-null   float64 
 27  rent_gt_20        11560 non-null   float64 
 28  rent_gt_25        11560 non-null   float64 
 29  rent_gt_30        11560 non-null   float64 
 30  rent_gt_35        11560 non-null   float64 
 31  rent_gt_40        11560 non-null   float64 
 32  rent_gt_50        11560 non-null   float64 
 33  universe_samples  11709 non-null   int64  
 34  used_samples      11709 non-null   int64  
 35  hi_mean           11587 non-null   float64 
 36  hi_median          11587 non-null   float64 
 37  hi_stdev           11587 non-null   float64
```

---

## Week 1, 2- Figure out the primary key and look for the requirement of indexing

---

```
In [14]: #UID is a unique userID value in the train and test dataset. So an index can be created from the UID feature
df_train.set_index(keys=['UID'],inplace=True)#Set the DataFrame index using existing columns.
df_test.set_index(keys=['UID'],inplace=True)

In [15]: df_train.head(2)
Out[15]:
   BLOCKID SUMLEVEL COUNTYID STATEID state state_ab city place type primary ... female_age_mean female_age_median female
   UID
267822  NaN    140      53     36  New York    NY Hamilton Hamilton City tract ... 44.48629 45.33333
246444  NaN    140     141     18  Indiana   IN South Bend Roseland City tract ... 36.48391 37.58333
2 rows × 79 columns
```

```
In [16]: df_test.head(2)
Out[16]:
   BLOCKID SUMLEVEL COUNTYID STATEID state state_ab city place type primary ... female_age_mean female_age_median female
   UID
255504  NaN    140      163     26  Michigan   MI Detroit Dearborn Heights CDP tract ... 34.78682 33.75000
282676  NaN    140      1     23  Maine     ME Auburn Auburn City City tract ... 44.23451 46.66667
2 rows × 79 columns
```

## Week 1, 3. Gauge the fill rate of the variables and devise plans for missing value treatment. Please explain explicitly the reason for the treatment chosen for each variable

---

```
In [17]: #percentage of missing values in train set
missing_list_train=df_train.isnull().sum()*100/len(df_train)
missing_values_df_train=pd.DataFrame(missing_list_train,columns=['Percentage of missing values'])
missing_values_df_train.sort_values(by=['Percentage of missing values'],inplace=True,ascending=False)
missing_values_df_train[missing_values_df_train['Percentage of missing values'] >0][10]
#BLOCKID can be dropped, since it is 100% missing values
```

```
Out[17]:
   Percentage of missing values
   BLOCKID          100.000000
   hc_samples        2.196113
   hc_mean           2.196113
   hc_median         2.196113
   hc_stddev         2.196113
   hc_sample_weight  2.196113
   hc_mortgage_mean 2.097288
   hc_mortgage_stddev 2.097288
   hc_mortgage_sample_weight 2.097288
   hc_mortgage_samples 2.097288
```

```
In [18]: #percentage of missing values in test set
missing_list_test=df_test.isnull().sum()*100/len(df_train)
missing_values_df_test=pd.DataFrame(missing_list_test,columns=['Percentage of missing values'])
missing_values_df_test.sort_values(by=['Percentage of missing values'],inplace=True,ascending=False)
missing_values_df_test[missing_values_df_test['Percentage of missing values'] >0][10]
#BLOCKID can be dropped, since it is 43% missing values
```

```
Out[18]:
   Percentage of missing values
   BLOCKID          42.857143
   hc_samples        1.061455
   hc_mean           1.061455
   hc_median         1.061455
   hc_stddev         1.061455
```

```
In [19]: df_train .drop(columns=['BLOCKID','SUMLEVEL'],inplace=True) #SUMLEVEL doest not have any predictive power and no variance
```

```
In [20]: df_test .drop(columns=['BLOCKID','SUMLEVEL'],inplace=True) #SUMLEVEL doest not have any predictive power
```

```
In [21]: # Imputing missing values with mean
missing_train_cols=[]
for col in df_train.columns:
    if df_train[col].isna().sum() !=0:
        missing_train_cols.append(col)
print(missing_train_cols)
```

```
[ 'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'r
```

```
In [22]: # Imputing missing values with mean
missing_test_cols=[]
for col in df_test.columns:
    if df_test[col].isna().sum() !=0:
        missing_test_cols.append(col)
print(missing_test_cols)
```

```
[ 'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'r
```

```
In [23]: # Missing cols are all numerical variables
for col in df_train.columns:
    if col in (missing_train_cols):
        df_train[col].replace(np.nan, df_train[col].mean(),inplace=True)
```

```
In [24]: # Missing cols are all numerical variables
for col in df_test.columns:
    if col in (missing_test_cols):
        df_test[col].replace(np.nan, df_test[col].mean(),inplace=True)
```

```
In [25]: df_train.isna().sum().sum()
Out[25]: 0
```

```
In [26]: df_test.isna().sum().sum()
Out[26]: 0
```

## Week 1, Exploratory Data Analysis (EDA):

Perform debt analysis. You may take the following steps: a) Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map. You may keep the upper limit for the percent of households with a second mortgage to 50 percent.

```
In [27]: from pandasql import sqldf
q1 = "select place,pct_own,second_mortgage,lat,lng from df_train where pct_own >0.10 and second_mortgage <0.5 order by second_mortgage desc limit 2500"
sqlDF = lambda q1: sqldf(q1, globals())
df_train_location_mort_pct=pysqldf(q1)
```

	place	pct_own	second_mortgage	lat	lng
0	Worcester City	0.20247	0.43363	42.254262	-71.800347
1	Harbor Hills	0.15618	0.31818	40.751809	-73.853582
2	Glen Burnie	0.22380	0.30212	39.127273	-76.835265
3	Egypt Lake-leto	0.11618	0.28972	28.029063	-82.495395
4	Lincolnwood	0.14228	0.28899	41.967289	-87.652434

```
In [28]: df_train_location_mort_pct.head()
```

```
Out[28]:
```

```
In [29]: import plotly.express as px
import plotly.graph_objects as go
```

```
In [30]: fig = go.Figure(data=go.Scattergeo(
    lat = df_train_location_mort_pct['lat'],
    lon = df_train_location_mort_pct['lon'],
))
fig.update_layout(
    geo=dict(
        scope = 'north america',
        showland = True,
        landcolor = "#rgb(212, 212, 212)",
        subunitcolor = "#rgb(255, 255, 255)",
        countrycolor = "#rgb(255, 255, 255)",
        showlakes = True,
        lakecolor = "#rgb(255, 255, 255)",
        showsubunits = True,
        showcountries = True,
        resolution = 50,
        projection = dict(
            type = 'conic conformal',
            rotation_ion = -108
        ),
        lonaxis = dict(
            showgrid = True,
            gridwidth = 0.5,
            range= [-140.0, -55.0 ],
            dtick = 5
        ),
        lataxis = dict (
            showgrid = True,
            gridwidth = 0.5,
            range= [ 20.0, 60.0 ],
            dtick = 5
        )
),
title='Top 2,500 locations with second mortgage is the highest and percent ownership is above 10 percent')
fig.show()
```


Top 2,500 locations with second mortgage is the highest and percent ownership is above 10 percent



b) Use the following bad debt equation: Bad Debt = P (Second Mortgage ∩ Home Equity Loan) Bad Debt = second\_mortgage + home\_equity - home\_equity\_second\_mortgage c) Create pie charts to show overall debt and bad debt

```
[31]: df_train['bad_debt']=df_train['second_mortgage']+df_train['home_equity']-df_train['home_equity_second_mortgage']
```

```
[32]: df_train['bins'] = pd.cut(df_train['bad_debt'],bins=[0,0.10,1], labels=["less than 50%","50-100%"])
df_train.groupby(['bins']).size().plot(kind='pie',subplots=True,startangle=90, autopct='%1.1f%%')
plt.axis('equal')
plt.show()
```



d) Create Box and whisker plot and analyze the distribution for 2nd mortgage, home equity, good debt, and bad debt for different cities

```
[33]: cols=[]
df_train.columns

[33]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
       'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',
       'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',
       'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
       'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',
       'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',
       'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
       'female_mean', 'female_median', 'female_stdev', 'female_sample_weight',
       'hc_mean', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
       'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
       'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
       'hs_degree_male', 'hs_degree_female', 'male_age_mean',
       'male_age_stdev', 'male_age_sample_weight', 'male_age_median',
       'male_age_samples', 'female_age_mean', 'female_age_stdev',
       'female_age_sample_weight', 'female_age_median',
       'pct_own', 'married', 'married_sep', 'separated', 'divorced',
       'bad_debt', 'bins'],
      dtype='object')
```


+ Code + Markdown

```
[34]: #Taking Hamilton and Manhattan cities data
cols=['second_mortgage','home_equity','debt','bad_debt']
df_box_hamilton=df_train.loc[df_train['city'] == 'Hamilton']
df_box_manhattan=df_train.loc[df_train['city'] == 'Manhattan']
df_box_city=pd.concat([df_box_hamilton,df_box_manhattan])
df_box_city.head(4)
```


COUNTYID	STATEID	state	state_ab	city	place	type	primary	zip_code	area_code	...	female_age_stdev	female_age_sample_weight	female_age_samples	pct_own	n	
267822	53	36	New York	NY	Hamilton	Hamilton	City	tract	13346	315	...	22.51276	685.33845	2618.0	0.79046	(4)
263797	21	34	New Jersey	NJ	Hamilton	Yardville	City	tract	8610	609	...	24.05831	732.58443	3124.0	0.64400	(4)
270979	17	39	Ohio	OH	Hamilton	Hamilton City	Village	tract	45015	513	...	22.66500	565.32725	2528.0	0.61278	(4)
259028	95	28	Mississippi	MS	Hamilton	Hamilton	CDP	tract	39746	662	...	22.79602	483.01311	1954.0	0.83241	(4)

4 rows × 79 columns


```
[35]: plt.figure(figsize=(10,5))
sns.boxplot(data=df_box_city,x='second_mortgage', y='city',width=0.5,palette="Set3")
plt.show()
```



```
[36]: plt.figure(figsize=(10,5))
sns.boxplot(data=df_box_city,x='home_equity', y='city',width=0.5,palette="Set3")
plt.show()
```




```
[37]: plt.figure(figsize=(10,5))
sns.boxplot(data=df_box_city,x='debt', y='city',width=0.5,palette="Set3")
plt.show()
```



[+ Code](#) [+ Markdown](#)

```
[38]: plt.figure(figsize=(10,5))
sns.boxplot(data=df_box_city,x='bad_debt', y='city',width=0.5,palette="Set3")
plt.show()
```

#Manhattan has higher metrics compared to Hamilton




[+ Code](#) [+ Markdown](#)


e) Create a collated income distribution chart for family income, house hold income, and remaining income

Please notice that the `distplot` is a deprecated function and will be removed in a future version. So, either use `displot` or `histplot`.

```
[39]: sns.distplot(df_train['hi_mean'])
plt.title('Household income distribution chart')
plt.show()
```




```
[40]: sns.distplot(df_train['family_mean'])
plt.title('Family income distribution chart')
plt.show()
```



```
[41]: sns.distplot(df_train['family_mean']-df_train['hi_mean'])
plt.title('Remaining income distribution chart')
plt.show()
#Income distribution almost has normality in its distribution
```


/opt/conda/lib/python3.7/site-packages/seaborn/distributions.py:2557: FutureWarning:  
'distplot' is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



## Week 2, Exploratory Data Analysis (EDA):


1- Perform EDA and come out with insights into population density and age. You may have to derive new fields (make sure to weight averages for accurate measurements):

```
[42]: fig, (ax1,ax2,ax3)=plt.subplots(3,1)
sns.distplot(df_train['pop'],ax=ax1)
sns.distplot(df_train['male_pop'],ax=ax2)
sns.distplot(df_train['female_pop'],ax=ax3)
plt.subplots_adjust(wspace=0.8,hspace=0.8)
plt.tight_layout()
plt.show()
```



[+ Code](#) [+ Markdown](#)

```
[43]: fig,(ax1,ax2)=plt.subplots(2,1)
sns.distplot(df_train['male_age_mean'],ax=ax1)
sns.distplot(df_train['female_age_mean'],ax=ax2)
plt.subplots_adjust(wspace=0.8,hspace=0.8)
plt.tight_layout()
plt.show()
```




[+ Code](#) [+ Markdown](#)

a) Use pop and ALand variables to create a new field called population density

```
[44]: df_train['pop_density']=df_train['pop']/df_train['ALand']
df_test['pop_density']=df_test['pop']/df_test['ALand']
```

```
[45]: sns.distplot(df_train['pop_density'])
plt.title('Population Density')
plt.show()
#Very less density is noticed
```




b) Use male\_age\_median, female\_age\_median, male\_pop, and female\_pop to create a new field called median age c) Visualize the findings using appropriate chart type

```
[46]: df_train['age_median']=(df_train['male_age_median']+df_train['female_age_median'])/2
df_test['age_median']=(df_test['male_age_median']+df_test['female_age_median'])/2
```


```
[47]: df_train[['male_age_median','female_age_median','male_pop','female_pop','age_median']].head()
```

UID	male_age_median	female_age_median	male_pop	female_pop	age_median
267822	44.00000	45.33333	2612	2618	44.666665
246444	32.00000	37.58333	1349	1284	34.791665
245683	40.83333	42.83333	3643	3238	41.833330
279653	48.91667	50.58333	1141	1559	49.750000
247218	22.41667	21.58333	2586	3051	22.000000

```
[48]: sns.distplot(df_train['age_median'])
plt.title('Median Age')
plt.show()
# Age of population is mostly between 20-60
# Majority are of age around 40
# Median age distribution has a gaussian distribution
# Some right skewness is noticed
```



```
[49]: sns.boxplot(df_train["age_median"])
plt.title("Population Density")
plt.show()
```



+ Code + Markdown

**Week 2, 2. Create bins for population into a new variable by selecting appropriate class interval so that the number of categories don't exceed 5 for the ease of analysis.**

```
[50]: df_train['pop'].describe()
```

```
[50]: count    27321.000000
count      4316.036685
std       2169.226173
min        0.000000
25%     2885.000000
50%     4042.000000
75%     5430.000000
max      53812.000000
Name: pop, dtype: float64
```

```
[51]: df_train['pop_bins']=pd.cut(df_train['pop'],bins=5,labels=['very low','low','medium','high','very high'])
```

```
[52]: df_train[['pop','pop_bins']]
```

```
[52]:   pop  pop_bins
```

UID	pop	pop_bins
267822	5230	very low
246444	2633	very low
245683	6881	very low
279653	2700	very low
247218	5637	very low
...	...	...
279212	1847	very low
277856	4155	very low
233000	2829	very low
287425	11542	low
265371	3726	very low

27321 rows × 2 columns

```
[53]: df_train['pop_bins'].value_counts()
```

```
[53]:
```

very low	27058
low	246
medium	9
high	7
very high	1

Name: pop\_bins, dtype: int64

a) Analyze the married, separated, and divorced population for these population brackets

```
[54]: df_train.groupby(by='pop_bins')[['married','separated','divorced']].count()
```

```
[54]:   married  separated  divorced
```

pop_bins	married	separated	divorced
very low	27058	27058	27058
low	246	246	246
medium	9	9	9
high	7	7	7
very high	1	1	1

```
[55]: df_train.groupby(by='pop_bins')[['married','separated','divorced']].agg(['mean', 'median'])  
#Very high population group has more married people and less percentage of separated and divorced couples  
#In very low population groups, there are more divorced people
```


```
[55]:
```

pop_bins	married		separated		divorced	
	mean	median	mean	median	mean	median
very low	0.507548	0.524680	0.019126	0.013650	0.100504	0.096020
low	0.584894	0.593135	0.015833	0.011195	0.075348	0.070045
medium	0.655737	0.618710	0.005003	0.004120	0.065927	0.064890
high	0.503359	0.335660	0.008141	0.002500	0.039030	0.010320
very high	0.734740	0.734740	0.004050	0.004050	0.030360	0.030360

+ Code + Markdown

b) Visualize using appropriate chart type

```
[56]: plt.figure(figsize=(10,5))  
pop_bin_married=df_train.groupby(by='pop_bins')[['married','separated','divorced']].agg(['mean'])  
pop_bin_married.plot(figsize=(10,8))  
plt.legend(loc='best')  
plt.show()
```



Week 2, 3. Please detail your observations for rent as a percentage of income at an overall level, and for different states.

```
[57]: rent_state_mean=df_train.groupby(by='state')['rent_mean'].agg(["mean"])
rent_state_mean.head()
```

	mean
	state
Alabama	774.004927
Alaska	1185.763570
Arizona	1097.753511
Arkansas	720.918575
California	1471.133857

```
[58]: income_state_mean=df_train.groupby(by='state')['family_mean'].agg(['mean'])
income_state_mean.head()
```

	mean
state	
Alabama	67030.064213
Alaska	92136.545109
Arizona	73328.238798
Arkansas	64765.377850
California	87655.470820

```
[59]: rent_perc_of_income=rent_state_mean['mean']/income_state_mean['mean']
rent_perc_of_income.head(10)
```

```
[59]: state
Alabama          0.011547
Alaska           0.012870
Arizona          0.014970
Arkansas         0.011131
California       0.016783
Colorado         0.013529
Connecticut      0.012637
Delaware          0.012929
District of Columbia 0.013198
Florida          0.015772
Name: mean, dtype: float64
```

```
[60]: #overall level rent as a percentage of income  
sum(df_train['rent_mean'])/sum(df_train['family_mean']))
```


[60]: 0.013358170721473864


Week 2, 4. Perform correlation analysis for all the relevant variables by creating a heatmap. Describe your findings.

```
[61]: df_train.columns
```

```
[62]: cor=df_train[['COUNTYID','STATEID','zip_code','type','pop','family_mean',
    'second_mortgage','home_equity','debt','hs_degree',
    'age_median','pct_own','married','separated','divorced']].corr()
```

```
[69]: plt.figure(figsize=(20,10))
sns.heatmap(cor,annot=True,cmap='coolwarm')
plt.show()
#High positive correlation is noticed between pop, male_pop and female_pop
#High positive correlation is noticed between rent_mean,h1_mean, family_mean, hc_mean
```





### Week 3, Data Pre-processing:

1. The economic multivariate data has a significant number of measured variables. The goal is to find where the measured variables depend on a number of smaller unobserved common factors or latent variables. 2. Each variable is assumed to be dependent upon a linear combination of the common factors, and the coefficients are known as loadings. Each measured variable also includes a component due to independent random variability, known as "specific variance" because it is specific to one variable. Obtain the common factors and then plot the loadings. Use factor analysis to find latent variables in our dataset and gain insight into the linear relationships in the data. Following are the list of latent variables:

- Highschool graduation rates
- Median population age
- Second mortgage statistics
- Percent own
- Bad debt expense

```
[65]: from sklearn.decomposition import FactorAnalysis
from factor_analyzer import FactorAnalyzer
```

```
[66]: fa=FactorAnalyzer(n_factors=5)
fa.fit_transform(df_train.select_dtypes(exclude= ('object','category')))
fa.loadings_
```

```
[66]: array([[-1.12589171e-01,  1.95646470e-02, -2.39331085e-02,
       -6.27612651e-02,  4.23474731e-02],
       [-1.10186764e-01,  1.33506211e-02,  2.79651256e-02,
       -1.49825865e-01,  1.10838806e-01],
       [-8.28678653e-02,  5.16372378e-02, -1.36451869e-01,
       -4.19181394e-02,  1.18038824e-01],
       [-1.80911394e-02,  1.92013789e-02,  5.81329950e-03,
       -1.64842763e-02, -6.12442386e-03],
       [ 9.92324766e-02, -9.72544292e-02, -6.546011353e-02,
       -1.33145930e-01, -1.48594597e-01],
       [-1.07335685e-02, -4.12376822e-02,  1.45853484e-01,
       8.80433198e-03,  1.88227568e-01],
       [-4.20796585e-02, -2.48107204e-02,  3.66726862e-02,
       -4.4594268e-02,  5.01180928e-02],
       [-2.44424306e-03, -1.53245418e-02, -2.68300860e-03,
       -4.52473040e-02,  2.37240553e-02],
       [ 7.92164302e-02,  9.57453324e-01, -8.71151615e-02,
       -6.59923641e-03, -3.97273239e-02],
       [ 7.39808109e-02,  9.18773520e-01, -1.88834837e-01,
       -2.79371152e-02, -3.54535320e-02],
       [ 8.0407879e-02,  9.47839224e-01, -6.08006491e-02,
       1.53627116e-02, -3.86877314e-02].
```

<< "if an error occur in FactorAnalysis that maybe the package doesn't downloaded. So, to download it: !pip install Factor\_Analyzer or pip install Factor\_Analyzer

### Week 4, Data Modeling :

1. Build a linear Regression model to predict the total monthly expenditure for home mortgages loan. Please refer 'deploment\_RX.xlsx'. Column hc\_mortgage\_mean is predicted variable. This is the mean monthly mortgage and owner costs of specified geographical location. Note: Exclude loans from prediction model which have NaN (Not a Number) values for hc\_mortgage\_mean.

```
[67]: df_train.columns
[67]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
       'primary', 'zip_code', 'area_code', 'lat', 'lng', 'Aland', 'AWater',
       'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',
       'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',
       'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
       'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',
       'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',
       'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
       'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
       'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
       'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
       'home_equity', 'home_mortgage', 'home_second_mortg', 'home_equity', 'debt',
       'second_mortg_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
       'hs_degree_male', 'hs_degree_female', 'male_age_mean',
       'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
       'male_age_samples', 'female_age_mean', 'female_age_median',
       'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
       'pct_own', 'married', 'married_snp', 'separated', 'divorced',
       'bad_debt', 'bins', 'pop_density', 'age_median', 'pop_bins'],
      dtype='object')
```

```
[68]: df_train['type'].unique()
type_dict={'type':{1,
                  'Urban':2,
                  'Town':3,
                  'CDP':4,
                  'Village':5,
                  'Borough':6}
                 }
df_train.replace(type_dict,inplace=True)
```

```
[69]: df_train['type'].unique()
```

```
[69]: array([1, 2, 3, 4, 5, 6])
```

```
[70]: df_test.replace(type_dict,inplace=True)
```

```
[71]: df_test['type'].unique()
```

```
[71]: array([4, 1, 6, 3, 5, 2])
+ Code + Markdown
```

```
[72]: feature_cols=['COUNTYID','STATEID','zip_code','type','pop', 'family_mean',
                  'second_mortgage', 'home_equity', 'debt','hs_degree',
                  'age_median','pct_own', 'married','separated', 'divorced']
```

```
[73]: x_train=df_train[feature_cols]
y_train=df_train['hc_mortgage_mean']
```

```
[74]: x_test=df_test[feature_cols]
y_test=df_test['hc_mortgage_mean']
```

```
[75]: from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error,mean_squared_error,accuracy_score
```

```
[76]: x_train.head()
```

	COUNTYID	STATEID	zip_code	type	pop	family_mean	second_mortgage	home_equity	debt	hs_degree	age_median	pct_own	married	separated	divorced
UID															
267822	53	36	13346	1	5230	67994.14790	0.02077	0.08919	0.52963	0.89288	44.666665	0.79046	0.57851	0.01240	0.08770
246444	141	18	46616	1	2633	50670.10337	0.02222	0.04274	0.60855	0.90487	34.791665	0.52483	0.34886	0.01426	0.09030
245683	63	18	46122	1	6881	95263.51431	0.00000	0.09512	0.73484	0.94288	41.833330	0.85331	0.64745	0.01607	0.10657
279653	127	72	927	2	2700	56401.68133	0.01086	0.01086	0.52714	0.91500	49.750000	0.65037	0.47257	0.02021	0.10106
247218	161	20	66502	1	5637	54053.42396	0.05426	0.05426	0.51938	1.00000	22.000000	0.13046	0.12356	0.00000	0.03109

```
[77]: sc=StandardScaler()
x_train_scaled=sc.fit_transform(x_train)
x_test_scaled=sc.fit_transform(x_test)
```

a) Run a model at a Nation level. If the accuracy levels and R square are not satisfactory proceed to below step.

```
[78]: linereg=LinearRegression()
linereg.fit(x_train_scaled,y_train)
```

```
[78]: LinearRegression()
```

```
[79]: y_pred=linereg.predict(x_test_scaled)
```

```
[80]: print("Overall R2 score of linear regression model", r2_score(y_test,y_pred))
print("Overall RMSE of linear regression model", np.sqrt(mean_squared_error(y_test,y_pred)))
#The Accuracy and R2 score are good, but still will investigate the model performance at state level
```

```
Overall R2 score of linear regression model 0.7348210754610929
Overall RMSE of linear regression model 323.1018894984635
```

b) Run another model at State level. There are 52 states in USA.

```
[81]: state=df_train['STATEID'].unique()
state[0:5]
#Picking a few IDs 20,1,45,6

[81]: array([36, 18, 72, 20, 1])

[82]: for i in [20,1,45]:
    print("State ID-",i)

    x_train_nation=df_train[df_train['COUNTYID']==i][feature_cols]
    y_train_nation=df_train[df_train['COUNTYID']==i]['hc_mortgage_mean']

    x_test_nation=df_test[df_test['COUNTYID']==i][feature_cols]
    y_test_nation=df_test[df_test['COUNTYID']==i]['hc_mortgage_mean']

    x_train_scaled_nation=sc.fit_transform(x_train_nation)
    x_test_scaled_nation=sc.fit_transform(x_test_nation)

    linereg.fit(x_train_scaled_nation,y_train_nation)
    y_pred_nation=linereg.predict(x_test_scaled_nation)

    print("Overall R2 score of linear regression model for state,"+str(i)+":",r2_score(y_test_nation,y_pred_nation))
    print("Overall RMSE of linear regression model for state,"+str(i)+":",np.sqrt(mean_squared_error(y_test_nation,y_pred_nation)))
    print("\n")
```

State ID- 20  
Overall R2 score of linear regression model for state, 20 :- 0.6046603766461811  
Overall RMSE of linear regression model for state, 20 :- 307.9718899931471

State ID- 1  
Overall R2 score of linear regression model for state, 1 :- 0.8104382475484617  
Overall RMSE of linear regression model for state, 1 :- 307.8275861848434

State ID- 45  
Overall R2 score of linear regression model for state, 45 :- 0.788744649785525  
Overall RMSE of linear regression model for state, 45 :- 225.69615420724125


+ Code + Markdown

```
[83]: #Check the residuals
residuals=y_test-y_pred
residuals
```

```
[83]: UID
230504 281.969088
252676 -69.935775
276314 190.761969
248614 -157.299627
286864 -9.887017
...
238088 -67.541646
242811 -41.955357
250127 -127.447569
241096 -330.820475
287763 217.760642
Name: hc_mortgage_mean, Length: 11780, dtypes: float64
```


```
[84]: plt.hist(residuals)
#Normal distribution of residuals
```

```
[84]: (array([6.000e+00, 3.000e+00, 2.900e+01, 7.670e+02, 7.823e+03, 2.716e+03,
       3.010e+02, 4.900e+01, 1.200e+01, 3.000e+00]),
 array([-2515.04284233, -1982.92661329, -1450.81038425, -918.69415521,
       -386.57792617, 145.53830267, 677.65453191, 1209.77076095,
       1741.88689999, 2274.00321903, 2806.11944807]),
```




```
[85]: sns.distplot(residuals)
```

```
[85]: <AxesSubplot:xlabel='hc_mortgage_mean', ylabel='Density'>
```



```
[86]: plt.scatter(residuals,y_pred)
# Same variance and residuals does not have correlation with predictor
# Independance of residuals
```

```
[86]: <matplotlib.collections.PathCollection at 0x7f97f45517d0>
```



**Please click on the project link above to see the output in more suitable way!**